



High Performance Computing and Rcpp

James Balamuta

Department of Informatics, Statistics
University of Illinois at Urbana-Champaign

July 24, 2017

CC BY-NC-SA 4.0, 2016 - 2017, James J Balamuta

On the Agenda

- High Performance Computing (HPC)
 - Processors
 - Performant code
- Compiled Code with Rcpp
 - C++ and Rcpp

On the Agenda

1 High Performance Computing

- Motivation
- Terminology
- Moving to HPC

2 Rcpp

- Motivation
- Setup

- Compilation
- File Naming
- Libraries
- C++ Code
- Addition Example
- Typing Issues
- Mean Function
- Proxy Model

Computing Quotation

“If you were plowing a field, which would you rather use? Two strong oxen or 1024 chickens?”
— Seymour Cray

Computing Quotation - Explained

"If you were plowing a field, which would you rather use? Two strong oxen or 1024 chickens?"

— Seymour Cray

- Lets assume for a moment that 2 oxen and 1024 chickens may provide the same amount of power.
- The question being asked is would you rather *manage* 2 operations or 1024 operations.
- In essence, 1024 would cause mass confusion as you may not be able to focus the each operation into usable power.

Power and Processors

When talking about computing, part of the **power** a computer has is given by the amount of information it can *process* within the Central Processing Unit.

Processing Lingo

Central Processing Unit (CPU):

A CPU is the brains of the computer that takes care of a majority of the calculations.





Processing Lingo Cont.

- **Core:**

- A core represents a single CPU.

- **Multi-core processor**

- A multi-core processor is a single component with two or more independent CPUs (cores) on the same die or block.



Processing Lingo Cont.

Thread:

A thread is a single line of commands that are getting processed by a CPU.

Mmmm... High Performance Computing

Definition:

High Performance Computing (HPC) most generally refers to the practice of aggregating computing power in a way that delivers much higher performance than one could get out of a typical desktop computer or workstation in order to solve large problems in science, engineering, or business.

Source

Myths, damn myths, and...

There are many myths that exist around HPC...

Here are some of them:

- Supercomputer or clusters only need apply
- Way too expensive (\$\$\$) to implement
- It's for Tech Firms or Academics
- Only useful for simulations
- There's no need for HPC in my field
- **This isn't available for R**

The Humor of the Situation

None of the myths are true!

In fact, we're barreling toward a future where it will be abnormal for code to run longer than 5 seconds.

Rear Admiral Grace Hopper and Nanoseconds



Figure 1: Grace Hopper, U.S. Navy, 1996 Vint Cerf

NCSA on why #HPCMatters!



Figure 9. NCSA's Leadership in #HPCMatters!

Why HPC? Why now?

- Part of the reason for **HPC** is the elephant in the room: **Big Data**.
- Computers have come a long way from requiring an entire room to do simple calculations to rendering movies and videos on a skinny jean pocket size iDevice.
- Though, to fully understand the *why* part, we need to talk about **Moore's law** and how it relates to computing.

Moore's Law

"The complexity. . . has increased roughly a factor of two per year. [It] can be expected to continue. . . for at least 10 years"

— Gordon Moore *"Cramming More Components onto Integrated Circuits," Electronics, pp. 114–117, April 19, 1965*


Commonly ***mis***stated as:

"Computer performance doubles every 18 months"


Video

Moore's Law - Transistors and Moores


Microprocessor Transistor Counts 1971-2011 & Moore's Law



Moore's Law - Clockspeeds of Processors



Moore's Law - Processors



Moore's Law and Reality




Figure 6: Mama Mia, that's a spicy meat-a-balla!

Full on Reality

- We are reaching / have reached a threshold with CPU performance.
- That is to say that we have become accustomed to PCs becoming faster every year but the future holds a different kind of **change**.
- The next saga is writing high performing code, improving compilers, and using parallelization.

On the Agenda

1 High Performance Computing

- Motivation
- Terminology
- Moving to HPC

2 Rcpp

- Motivation
- Setup

- Compilation
- File Naming
- Libraries
- C++ Code
- Addition Example
- Typing Issues
- Mean Function
- Proxy Model

R is dead, long live *R*?

- *R* is a wonderful language that allows for just about **anything**
- Each time *R* dies, it comes back stronger courtesy of community involvement.
- Every weakness gets addressed and extensions are made.

The Problem with *R*

"I like to think of R as one of the best programming languages with one of the worst 'standard libraries' "

— Kevin Ushey in *Needless Inefficiencies in R – Head and Tail*

- The downsides of a language for *statisticians by statisticians*.

R is dead, long live R ?

The main staying power of R is:

- Interactively work with code
- Rapid prototyping (thanks to the interpreter)
- Syntax for statisticians by statisticians to explore data
- Top notch statistical methods
- Easy install and distribution (thanks [CRAN!](#))
- Gateway or interface to other applications.

At the same time, R is incredibly weak in:

- Speed (downside of prototyping)
- Loops (*shudders*)
- Effective memory management (a bit aggressive on allocations)
- Multicore support (default compile of R 's BLAS is single core!)

Why compile code for R?

The why for compiling code is able to be reduced to:

- Speed
 - Making things go fast is fun and beneficial in this age of instantaneous response.
- More libraries and tools
 - Libraries outside of the scope of R (boost, GSL, Eigen, Armadillo) are now able to be used!
- Great Support
 - There are countless answers to C++ questions and people willing to answer them! e.g. stackoverflow.com
- C++ is rock solid
 - C++ has been battle tested in industry and academia alike.

Why *Rcpp*?

Rcpp is:

- **Well Documented**

- Rcpp Gallery has 100+ posts and over 1100 unit tests

- **Straightforward to Use (with RStudio)**

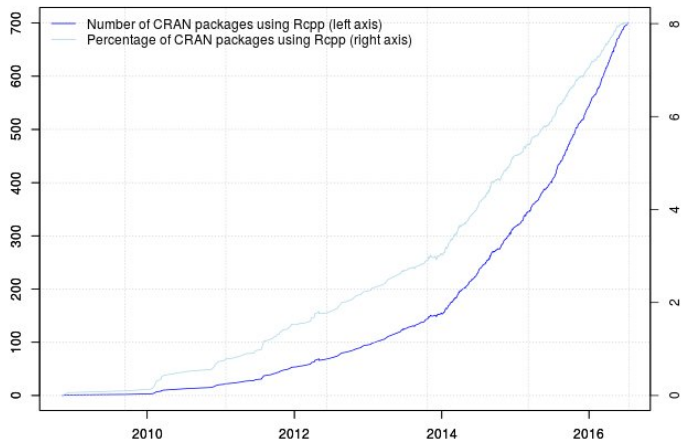
- Press a button and code compiles! No worries about using terminal.

- **Seamless access to all R objects**

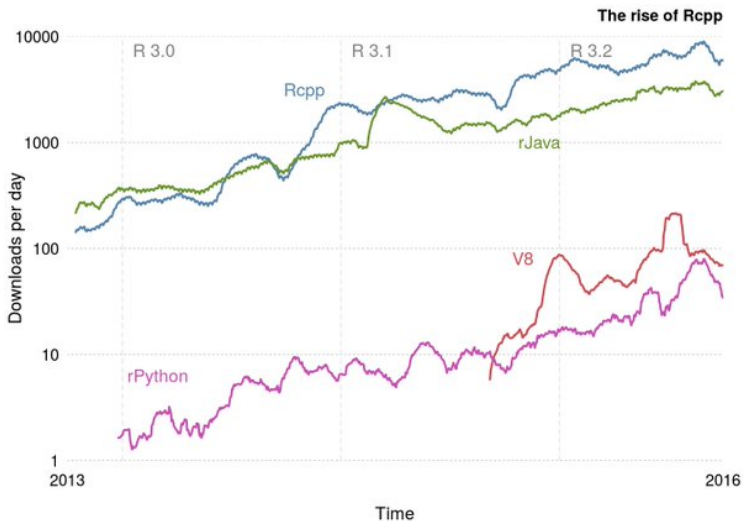
- Move back and forth between *R* and *C++* with easy.

Rise of Rcpp

Growth of Rcpp usage on CRAN



Rise of Rcpp



C++ Disclaimer

- C++ is a very **powerful** language.
- We're only going to briefly cover parts that relate to *non*-complex operations.
 - See CS 225 for full treatment
- E.g things that *R* is inherently bad at like loops. . .

Setting Up the Environment

Depending on the *type* of Operating System, there are different types of installations required.


Please see the guide for each operating system:

- [Windows](#) (\geq Windows 7)
- [macOS \(basic\)](#) **or** [parallel on macOS \(advanced\)](#)


Creating C++ Files in RStudio

There are two ways to create a Cpp file within RStudio:

- **File ⇒ New File ⇒ C++ File**

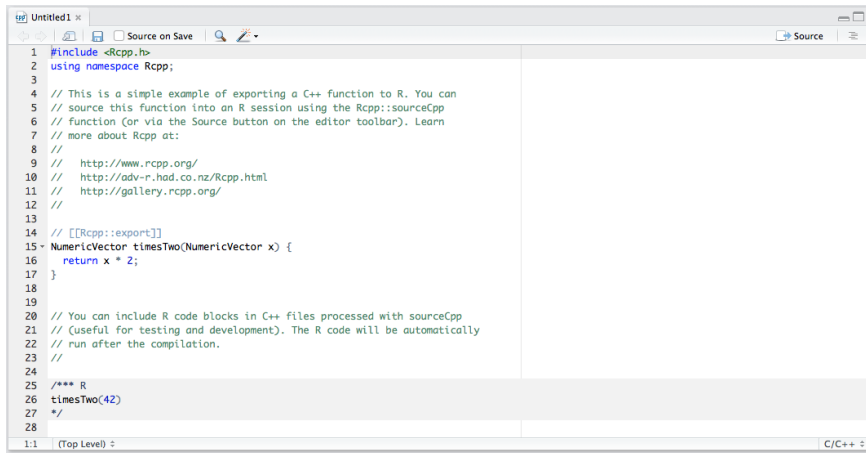


- **New Document Symbol ⇒ C++ File**



Results In...

- Standard Rcpp templated file with an *example*



```
1 #include <Rcpp.h>
2 using namespace Rcpp;
3
4 // This is a simple example of exporting a C++ function to R. You can
5 // source this function into an R session using the Rcpp::sourceCpp
6 // function (or via the Source button on the editor toolbar). Learn
7 // more about Rcpp at:
8 //
9 // http://www.rcpp.org/
10 // http://adv-r.had.co.nz/Rcpp.html
11 // http://gallery.rcpp.org/
12 //
13
14 // [[Rcpp::export]]
15 NumericVector timesTwo(NumericVector x) {
16   return x * 2;
17 }
18
19
20 // You can include R code blocks in C++ files processed with sourceCpp
21 // (useful for testing and development). The R code will be automatically
22 // run after the compilation.
23 //
24
25 /** R
26 timesTwo(42)
27 */
28
```


Compiling...

Three different ways to trigger a compilation of C++ code in a .cpp file:

- Use one of RStudio's keyboard shortcuts!
 - All OSes: Ctrl + Shift + Enter
- Type into console:

```
Rcpp::sourceCpp("rcpp_twotimes.cpp")
```

- Press the Source button at the top right of the editor window.



Output ...

```
timesTwo(42)  
## [1] 84
```

Note: This output was **automatic** after compiling the code due to the comment at the end.

Kinds of Compiling Techniques with Rcpp

There are three ways to compile code without embedding it within an R package.

1 evalCpp()

- To quickly check different C++ expressions.
- Limited scope

2 cppFunction()

- For defining inline function code
- Limit 1 function per call.

3 sourceCpp()

- For code kept in an alternative file
- Multiple files with interfunction dependence.
- This is the **preferred** way to work with Rcpp code.

Main compilation technique

- For all intents and purposes, we will use `sourceCpp()` to compile C++ code.
- `sourceCpp()` is better as it allows for C++ syntax highlighting whereas `cppFunction()` loses the C++ syntax highlighting due to the string context.

C++ File Names

- When working with C++, avoid using spaces or special symbols in either the file path or file name.
 - File Name: `example.cpp`
 - File Path: `/home/netid/example.cpp`
- If you need to use a space, use the underscore: `_`

Examples:

- Good: `rcpp_example.cpp`, `hello2rcpp.cpp`
- Bad: `C++ Example.cpp`, `rcpp is the bees knees.cpp`

C++ vs. R - Libraries

- To include different libraries using header (.h) files akin to R's `library()` function, write in the C++ file:

```
#include <Rcpp.h>           // Includes the Rcpp C++ header  
                             // Akin to calling library(Rcpp)  
                             // in C++
```

- Unlike R, we also have to explicitly add the namespace of a header that we wish to use.

```
using namespace Rcpp; // C++ search scope
```

C++ vs. R - Libraries

All C++ files must therefore have at the top of them:

```
#include <Rcpp.h>      // Includes the Rcpp C++ header  
using namespace Rcpp; // C++ search scope
```


C++ vs. R - Note on Namespaces of Libraries

- Specifying the namespace avoids having to prefix function calls with the `Rcpp::` namespace.
- **Note:** This is *not* a good style but it makes beginning in C++ a bit easier.
- In particular, ambiguity is introduced into the code when two libraries provide functions with the same name (e.g. `std::sqrt` and `arma::sqrt`).
- *R* would warn when this overload happens on package load, *C++* will **not**.

C++ vs. R - Commenting in Code

Comments in C++ come in two flavors:

```
/* Group comment  
   Across Multiple Lines  
   */  
  
// Single line comment
```

The use of R's traditional comment:

```
# pound/hash comment
```

is used to declare preprocessor macros and, thus, should **not** be used within a C++ file.

Steps toward a C++ Function

- Before writing a C++, try to write the function in *R* first.
- Consider an *R* function called `hello` whose goal is to print “Hello R/C++ World!”

```
hello = function(){  
  cat("Hello R/C++ World!\n")  
}
```

C++ vs. R - First Function in C++

We can mimic this by creating the following C++ function.

```
// Akin to cat("Hello R/C++ World!\n")  
void hello() { // Declaration  
  Rcout << "Hello R/C++ World!\n"; // cat() in C++  
}
```

Differences between *R* and C++:

- Return type of the function `void` is specified before function name.
- `void` indicates that no information is *returned*.
- `cat()` is done using `Rcout` with strings being delimiters of by `<<`.

First Function in C++ within R

In order to write a C++ function that works with *R*, you must specify the intent to export into *R* by writing directly above the function the **Rcpp Export Attribute**

```
// [[Rcpp::export]]
```

e.g.

```
#include <Rcpp.h>
using namespace Rcpp;

// [[Rcpp::export]]
void hello() {
  Rcout << "Hello R/C++ World!\n";
}
```

Compile it!

Pick one of the ways listed previously to compile:

```
#include <Rcpp.h>
using namespace Rcpp;

// [[Rcpp::export]]
void hello() {
  Rcout << "Hello R/C++ World!\n";
}
```

Calling the C++ Function within R

C++ functions are automatically surfaced into the *R* environment by *Rcpp* under their defined name. So, in the previous example, we would have `hello()` in the global environment that we can now call.

```
# Call C++ Code like a normal R function  
hello()
```

```
## Hello R/C++ World!
```

Automatically Run R Code on C++ Compile

To automatically test code after compile, you can embed the R code in special C++ group comment blocks like so:

```
#include <Rcpp.h>
using namespace Rcpp;

// [[Rcpp::export]]
void hello() {
  Rcout << "Hello R/C++ World!\n";
}

/** R
# This is R code in the C++ code file!
hello()
*/
```


C++ Function within *R*

The insides of the C++ function are slightly different

```
hello
```

```
## function ()
```

```
## invisible(.Primitive(".Call")(<pointer: 0x1079f3930>))
```

This is only problematic when you are trying to do cluster computing (more later).

Parameter Functions in R

- The beauty of a function is being able to slightly change variables and obtain a new input.
- To add only *two* numbers, an `add()` can be created as:

```
add = function(a,b){  
  return(a + b)  
}
```

```
add(0L, 2L)    # Remember L means integer!
```

```
## [1] 2
```

```
add(2.5, 1.1) # Double/numeric
```

```
## [1] 3.6
```

C++ vs. R - Parameter Functions

Difference between an *R* function and that of a *C++* function are as follows:

- 1 Return data type
- 2 Data type of input parameters

```
#include <Rcpp.h>
using namespace Rcpp;           // Import Statements

// [[Rcpp::export]]
double addRcpp(double a, double b) { // Declaration
    double out = a + b; // Add `a` to `b`
    return out;         // Return output
}
```

Note: Everything must be **pre-typed**.

Calling the C++ `addRcpp()` function

- You may have noticed that there should be a *typing* issue that arises.
- *Rcpp* is kind and allows for the seamless conversion of integers in *R* to doubles in C++.

```
add(0L, 2L)    # Integers into double
```

```
## [1] 2
```

```
add(2.5, 1.1) # Double into double
```

```
## [1] 3.6
```

Calling another C++ function with different types

```
#include <Rcpp.h>
using namespace Rcpp;           // Import Statements

// [[Rcpp::export]]
double addRcpp(double a, double b) { // Declaration
    double out = a + b;             // Add `a` to `b`
    return out;                     // Return output
}

// [[Rcpp::export]]
int addRcppInt(int a, int b) { // Declaration
    return addRcpp(a, b);        // Call previous function
}
```

Calling the C++ `addRcpp()` function with clashing types

Note, C++ will even *try* to handle the conversion between `int` and `double`.

```
addRcppInt(2.5, 1.1)  # Call in *R*
```

```
## [1] 3
```

Follows *bias* rounding procedure of:

- If $x \leq 0.5$, then round down: $\lfloor y \rfloor$ (floor)
- If $x > 0.5$, then round up: $\lceil y \rceil$ (ceiling)

where x is the fractional component of z and y is the integer component of z , e.g. $x = 0.3, y = 2 \Rightarrow z = 2.3$

- Why is the output 3 instead of 4?

C++ vs. R - Functions

Data Types to choose from:

- `double`: 1 numeric
- `int`: 1 integer
- `std::string`: 1 character
- `void`: nothing
- `*Vector`: vector of either Integer, Numeric, or Character
- `*Matrix`: matrix of either Integer, Numeric, or Character

Pay Attention to Your Data Types!

R - Mean Function

Recall that

$$\bar{x}_n = \frac{1}{n} \sum_{i=1}^n x_i$$

Given by:

```
muR = function(x) {  
  sum_r = 0  
  for (i in seq_along(x)) {  
    sum_r = sum_r + x[i]  
  }  
  
  sum_r / length(x)  
}
```


C++ - Mean Function

Goal: Obtain the Mean of a vector.

```
#include <Rcpp.h>
using namespace Rcpp;           // Import Statements

// [[Rcpp::export]]
double muRcpp(NumericVector x) { // Declaration

    int n = x.size();           // Find the vector length
    double sum_x = 0;           // Set up storage

    for(int i = 0; i < n; ++i) { // For Loop in C++
        // Shorthand for sum_x = sum_x + x[i];
        sum_x += x[i];
    }

    return sum_x / n;           // Return division
}
```

Checking C++ and R Function Equality

Check the equality of both functions using `all.equal()`

```
# Done in *R*
```

```
set.seed(112) # Set seed
```

```
x = rnorm(10) # Generate data
```

```
all.equal(muRcpp(x), muR(x)) # Test Functions
```

```
## [1] TRUE
```

Rcpp Proxy Model

- All objects using the `*Vector` or `*Matrix` tags are **proxy** objects.
- That means, they are acting as pointers to the actual memory.
- So, the call to function has *Rcpp* objects being passed by **reference**.

Rcpp Proxy Model - Example

```
#include <Rcpp.h>
using namespace Rcpp;           // Import Statements

// [[Rcpp::export]]
void ref_ex(IntegerVector x) { // Declaration
    x = x + 1; // Add 1 and save it to x via Rcpp Sugar
}
```

```
(x = 1:10) # Span from 1 to 10
```

```
## [1] 1 2 3 4 5 6 7 8 9 10
```

```
ref_ex(x) # No output due to no return
x         # Different Span!
```

```
## [1] 2 3 4 5 6 7 8 9 10 11
```

Rcpp Proxy Model - Trouble Ho!

The proxy model works reasonably well to a degree...

```
(x = seq(0.5,5.5))  # Span from 0.5 to 5.5
```

```
## [1] 0.5 1.5 2.5 3.5 4.5 5.5
```

```
ref_ex(x)           # No output due to no return  
x                   # Same span????
```

```
## [1] 0.5 1.5 2.5 3.5 4.5 5.5
```

Rcpp Proxy Model - The failures

- Note that the type of `x` is `numeric` in this case and not `integer` like before.
 - Check yourself via:

```
typeof(x)
```

- Thus, the Rcpp proxy model **failed** to use existing memory and created a new object since it expected type `IntegerVector` but had to cast from `NumericVector`.

Lesson of the Day...

Pay Attention to Your Data Types!

Misc Note on Compiling Techniques with Rcpp

- When calling both the `cppFunction()` and `evalCpp()` there is a call to the `sourceCpp()` to generate the object. This is a really effective code design.
- For details see: [Rcpp Attributes Vignette](#) and [Code Source](#)